



國立臺灣大學

National
Taiwan
University

Lab2: Threads

EE 5173 Operating System

TA: Hsueh-Yi, Chen. Advisor: Farn, Wang.

2025/4/9

Lab2 Description

- Minimally functional thread system has been provided. Your job is to extend the functionality of this system to gain a better understanding of synchronization problems.
- Highly recommend to read the reference. (We just selected some content here.)
https://web.stanford.edu/class/cs140/projects/pintos/pintos_2.html#SEC15

Strongly Recommended To Read

- A.1 Loading
- A.2 Threads
- A.3 Synchronization
- A.4 Interrupt Handling
- A.5 Memory Allocation

https://web.stanford.edu/class/cs140/projects/pintos/pintos_6.html#SEC91

Background

Understanding Threads

- When a thread is created, you are creating a new context to be scheduled. You provide a function to be run in this context as an argument to **thread_create()**.

```
/* Creates a new kernel thread named NAME with the given initial
   * PRIORITY, which executes FUNCTION passing AUX as the argument,
   * and adds it to the ready queue. Returns the thread identifier
   * for the new thread, or TID_ERROR if creation fails.
   */
```

```

   If thread_start() has been called, then the new thread may be
   scheduled before thread_create() returns. It could even exit
   before thread_create() returns. Contrariwise, the original
   thread may run for any amount of time before the new thread is
   scheduled. Use a semaphore or some other form of
   synchronization if you need to ensure ordering.

```

```

   The code provided sets the new thread's `priority' member to
   PRIORITY, but no actual priority scheduling is implemented.
   Priority scheduling is the goal of Problem 1-3. */

```

```

tid_t
thread_create (const char *name, int priority,
              thread_func *function, void *aux)
{
    struct thread *t;
```

Understanding Threads

- At any given time, exactly one thread runs and the rest, if any, become inactive.
- The scheduler decides which thread to run next. (If no thread is ready to run at any given time, then the special "idle" thread, implemented in `idle()`, runs.)
- Synchronization primitives can force context switches when one thread needs to wait for another thread to do something.
- The mechanics of a context switch are in `threads/switch.S`

Understanding Threads

- It is recommended to use a GDB debugger and printf to dynamically trace through a context switch.
- Set a breakpoint on **schedule()** to start out, and single step from there.

Source Files

- threads
- devices
- lib

https://web.stanford.edu/class/cs140/projects/pintos/pintos_2.html#SEC15

Synchronization

1. Synchronization in Pintos

- Proper synchronization is crucial for solving concurrency problems.
- Disable interrupts to solve synchronization problems is tempting but should be avoided.

2. Disabling Interrupts for Synchronization

- Disabling interrupts is useful only for coordinating data between kernel threads and interrupt handlers.
- Interrupt handlers cannot sleep, so they can't acquire locks.
- Data shared between kernel threads and interrupt handlers must be protected by turning off interrupts.

Synchronization

3. Only requires accessing a little bit of thread state from interrupt handlers

- For the alarm clock, the timer interrupt needs to wake up sleeping threads
- In the advanced scheduler, the timer interrupt needs to access a few global and per-thread variables

Disabling interrupts ensures that the timer interrupt does not interfere.

Synchronization

4. Here are some reminders about disabling interrupts

- When disabling interrupts, minimize the code section to avoid losing important events like timer ticks or input.
- Excessive interrupt disabling increases latency, which can make the system feel sluggish.

5. Others

- The synchronization primitives themselves in `synch.c` are implemented by disabling interrupts. You may need to increase the amount of code that runs with interrupts disabled here, but you should still try to keep it to a minimum.
- Disabling interrupts can be useful for debugging, if you want to make sure that a section of code is not interrupted

LAB

Alarm Clock

- Reimplement the `timer_sleep()` function (defined in `devices/timer.c`) to avoid busy waiting.
- Function: `void timer_sleep (int64_t ticks)`
- Hint: We need a place to store (queue) information about sleeping threads and wake them up by moving them to the ready queue

```
void
timer_sleep (int64_t ticks)
{
    int64_t start = timer_ticks ();

    ASSERT (intr_get_level () == INTR_ON);
    while (timer_elapsed (start) < ticks)
        thread_yield ();
}
```

Priority Scheduling

- When a thread is **added to the ready list that has a higher priority** than the currently running thread, the current thread should **immediately yield** the processor to the new thread.
- Similarly, when threads are waiting for a lock, semaphore, or condition variable, the highest priority waiting thread should be **awakened first**.
- A thread may raise or lower its own priority at any time, but lowering its priority such that it no longer has the highest priority must cause it to **immediately yield the CPU**.
- PRI_MIN (0) to PRI_MAX (63), from lowest to highest

Priority Scheduling (priority donation)

- One issue with priority scheduling is "priority inversion".
- Consider high, medium, and low priority threads H, M, and L, respectively. If H needs to wait for L (for instance, for a lock held by L), and M is on the ready list, then H will never get the CPU because the low priority thread will not get any CPU time.
- A partial fix for this problem is for H to "donate" its priority to L while L is holding the lock, then recall the donation once L releases (and thus H acquires) the lock.

Priority Scheduling (priority donation)

- You will need to account for all different situations in which priority donation is required. (Test Driven)
- Be sure to handle multiple donations, in which multiple priorities are donated to a single thread. You must also handle nested donation: if H is waiting on a lock that M holds and M is waiting on a lock that L holds, then both M and L should be boosted to H's priority. (Recursive call)

Priority Scheduling (priority donation)

threads/thread.c

Function: void thread_set_priority(int new_priority)

- Sets the current thread's priority to new_priority. If the current thread no longer has the highest priority, yields.

Function: int thread_get_priority(void)

- Returns the current thread's priority. In the presence of priority donation, returns the higher (donated) priority.

~~Threads modify other threads' priorities.~~

Advanced Scheduler

- Implement a multilevel feedback queue scheduler similar to the BSD scheduler to **reduce the average response time** for running jobs on your system. BSD Scheduler
- Recommend complete priority scheduler first before start work on it (Advanced scheduler choose the thread to run based on priorities).
- The BSD scheduler is enabled, threads no longer directly control their own priorities

https://web.stanford.edu/class/cs140/projects/pintos/pintos_7.html#SEC131

Advanced Scheduler

- $\text{priority} = \text{PRI_MAX} - (\text{recent_cpu} / 4) - (\text{nice} * 2)$
- $\text{recent_cpu} = (2 * \text{load_avg}) / (2 * \text{load_avg} + 1) * \text{recent_cpu} + \text{nice}$
- $\text{load_avg} = (59/60) * \text{load_avg} + (1/60) * \text{ready_threads}$

Advanced Scheduler

- You can add a file called “fixed-point.h”
- Recommendation: inline static function

```
devices/timer.c      | 42 +++++-
threads/fixed-point.h | 120 ++++++
threads/synch.c      | 88 ++++++
threads/thread.c     | 196 ++++++
threads/thread.h     | 23 +++
5 files changed, 440 insertions(+), 29 deletions(-)
```

Convert n to fixed point:	$n * f$
Convert x to integer (rounding toward zero):	x / f
Convert x to integer (rounding to nearest):	$(x + f / 2) / f$ if $x \geq 0$, $(x - f / 2) / f$ if $x < 0$.
Add x and y :	$x + y$
Subtract y from x :	$x - y$
Add x and n :	$x + n * f$
Subtract n from x :	$x - n * f$
Multiply x by y :	$((\text{int64_t}) x) * y / f$
Multiply x by n :	$x * n$
Divide x by y :	$((\text{int64_t}) x) * f / y$
Divide x by n :	x / n

Grading Policy

- Total - 100%
 - Test Score – 40%
 - Tech Report – 60%

Deadline 2025/4/30 23:59

- Filename (-3%)
- Late submission (-10%), over 1 weeks (-20%) over 2 weeks (Immediate 0 points)
- Plagiarism or All/Most LLM (Immediate 0 points)
- Bonus (Other implementations: Max 10)

Report

- Content format: 12pt front, 16pt row height, and align to the left.
- Caption format: 18pt and Bold font.
- Figure: center with single line row height.
- Upload with the file structure format :

G[team number]_2.zip (e.g. G01_2.pdf)

| — G[team number]_2.pdf

| — pintos

| └ devices

| └ ...

|

Requirement in tech report

- Please follow the design document questions

(<https://web.stanford.edu/class/cs140/projects/pintos/threads.tmpl>)

- **Test Results**

- Provide a snapshot of test results and your **make check** score.
- If any tests fail, briefly explain the possible reasons.

```
pass tests/threads/mlfqs-load-60
pass tests/threads/mlfqs-load-avg
pass tests/threads/mlfqs-recent-1
pass tests/threads/mlfqs-fair-2
pass tests/threads/mlfqs-fair-20
pass tests/threads/mlfqs-nice-2
pass tests/threads/mlfqs-nice-10
pass tests/threads/mlfqs-block
All 27 tests passed.
```



Thanks for your attention !

